

CSCI235 Database Systems

MongoDB Aggregation Framework

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

2/35

Aggregation framework

Aggregation framework is a query language that that can be used to **transform** and to **combine** data from multiple documents in order to **generate** new information not available in any single document

Aggregation framework in MongoDB is similar to **SQL GROUP BY** clause and **HAVING** clause

Aggregation framework makes a task database search much easier and more efficient through specification of a series of operations in an array and processing it in a single call

Aggregation framework defines an **aggregation pipeline** where the output from each step in the pipeline provides input to the next step

Every step in a **pipeline** executes a single operation on the **input documents** to transform the **input** and to generate **output document**

A **pipeline** processes a **stream of documents** through several operations like **filtering**, **projecting**, **grouping**, **sorting**, **limiting**, and **skipping**

The same operations can be repeated many times in a **pipeline** in any order

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

3/35

Aggregation framework

Some operations that can be used in a **pipeline**:

- **\$project**: Extracts the components of a documents to be placed in an output document (similar to **SELECT** clause)
- **\$match**: Filters the documents to be processed, similar to **find()** (and similar to **WHERE** clause)
- **\$limit**: Limits the total number of documents to be passed to the next operation (similar to **LIMIT** clause, **rownum** condition)
- **\$skip**: Skips a given specified number of documents
- **\$unwind**: Expands (unnest) an array, generating one output document for each array entry
- **\$group**: Groups documents by a specified key (**GROUP BY** clause)
- **\$sort**: Sorts documents (**ORDER BY** clause)
- **\$out**: Saves the results from a **pipeline** to a collection
- **\$count**: Counts the total number of documents in a **pipeline**

Sample collection

A sample document

```
db.department.insert(  
  { "name": "School of Astronomy",  
    "code": "SOA",  
    "total_staff_number": 25,  
    "budget": 10000,  
    "address": { "street": "Franz Josef Str",  
                  "bldg": 4,  
                  "city": "Vienna",  
                  "country": "Austria" },  
    "courses": [ { "code": "SOA101",  
                    "title": "Astronomy for Kids",  
                    "credits": 3 },  
                  { "code": "SOA201",  
                    "title": "Black Holes",  
                    "credits": 6 },  
                  { "code": "SOA301",  
                    "title": "Dark Matter",  
                    "credits": 12 }  
                ]  
  }  
);
```

Sample document

Aggregations

Outline

Aggregation framework

Operations

\$project

\$match

\$limit

\$skip

\$unwind

\$group

\$sort

\$out

\$count

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

6/35

\$project

\$project extracts components of subdocuments, rename components, and performs operations on components

Select name of each department, skip document identifier

```
db.department.aggregate([ {$project:{"name":1,"_id":0}} ])
```

aggregate()

Select name of each department and rename name component

```
db.document.aggregate([ {$project:{"Department name":"$name"}} ])
```

aggregate()

"\$keyname" syntax is used to refer to a value associated with a key **"keyname"** in the aggregation framework

Select a name of each department and concatenate it with its code

```
db.document.aggregate([ {$project:{"Name and code":  
    {"$concat":["$name","-","$code"]}}} ])
```

aggregate()

\$project

Select a name of department and 10% of its budget

```
db.department.aggregate([  { $project: { "name": 1,
                                         "10% of budget": { "$multiply": [ "$budget", 0.1 ] },
                                         "_id": 0 }
                           ]
                           )
```

aggregate()

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

9/35

\$match

Operation `$match` selects the documents that satisfy a given condition

Find the departments with `budget > 10000`

```
db.department.aggregate([ {$match:{"budget":{$gt:10000}}} ])
```

`aggregate()`

Find the `names` and `codes` of departments with `budget > 10000`

```
db.department.aggregate([ {$match:{"budget":{$gt:10000}}},  
                          {$project:{"name":1,"code":1,"_id":0}} ])
```

`aggregate()`

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

11/35

\$limit

Operation `$limit` passes a given number of documents through a pipeline

Find the first 2 documents

```
db.department.aggregate([ {$limit:2} ])
```

`aggregate()`

Find the first document with `budget > 1000`

```
db.department.aggregate([ {$match:{"budget":{"$gt":1000}}},  
                           {$limit:1} ])
```

`aggregate()`

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

13/35

\$skip

Operation **\$skip** eliminates a given number of documents from a pipeline

List all documents in a collection except the first two

```
db.department.aggregate([ {$skip:2} ])
```

aggregate()

List all documents with a **budget** greater than 10000 except the first one

```
db.department.aggregate([ {$match:{"budget":{"$gt":10000}}},  
                           {$skip:1}] )
```

aggregate()

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

15/35

\$unwind

Operation `$unwind` creates a separate document for each element of a given array

A document is replicated for each element of an array, i.e. an array is **unnested**

For each department and for each course offered by a department create a separate document

```
db.department.aggregate([ {$unwind:"$courses"} ])
```

`aggregate()`

For each department and for each course offered by a department create a separate document, list only the courses

```
db.department.aggregate([ {$unwind:"$courses"},  
                           {$project:{"courses":1,"_id":0}} ])
```

`aggregate()`

\$unwind

Use **aggregation framework** and **\$unwind** operation to list the codes of all courses

```
db.department.aggregate([ {$unwind: "$courses"},  
                          {$project: {"code": "$courses.code", "_id": 0}} ])
```

aggregate()

```
{ "code" : "CSCI835" }  
{ "code" : "CSIT115" }  
{ "code" : "CSCI317" }  
{ "code" : "CSIT321" }  
{ "code" : "SOA101" }  
{ "code" : "SOA201" }  
{ "code" : "SOA301" }  
{ "code" : "SOPH101" }  
{ "code" : "SOPH102" }  
{ "code" : "SOPH103" }
```

Results

\$unwind

Use **aggregation framework** and **\$unwind** operation to find all courses with 12 credits

```
db.department.aggregate([ {$unwind:"$courses"},  
                          {$project:{"courses":1,"_id":0}},  
                          {$match:{"courses.credits":12}} ])
```

aggregate()

Aggregations

Outline

Aggregation framework

Operations

\$project

\$match

\$limit

\$skip

\$unwind

\$group

\$sort

\$out

\$count

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

19/35

\$group

Operation `$group` groups the documents and applies the aggregation functions to each group

List non-distinct values of `total_staff_number`

```
db.department.aggregate([{$project:{"total_staff_number":1,"_id":0}}])
```

aggregate()

```
{ "total_staff_number" : 25 }  
{ "total_staff_number" : 5 }  
{ "total_staff_number" : 25 }
```

Results

Group the documents by `total_staff_number` and list the distinct values of `total_staff_number`

```
db.department.aggregate([{$group:{"_id":"$total_staff_number"}}])
```

aggregate()

```
{ "_id" : 5 }  
{ "_id" : 25 }
```

Results

\$group

Perform groupings as above and rename `_id` to `total_staff_number`

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number"}},  
                          {$project:{"total_staff_number":"$id","_id":0}} ])
```

aggregate()

```
{ "total_staff_number" : 5 }  
{ "total_staff_number" : 25 }
```

Results

Find the total number of distinct values of `total_staff_number`

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number"}},  
                          {$count:"Total distinct values"}])
```

aggregate()

Find the total number of distinct values of `total_staff_number`

```
db.department.distinct("total_staff_number").length
```

distinct()

\$group

Group the documents by `total_staff_number` and by `budget`

```
db.department.aggregate([ {$group:{"_id":{"total_staff_number":"$total_staff_number",  
                                     "budget":"$budget"}}}  ])
```

Group the documents by `total_staff_number` and count the total number of departments in each group

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                     "total_departments":{"$sum:1"}}}  ])
```

```
{ "_id" : 5, "total_departments" : 1 }  
{ "_id" : 25, "total_departments" : 2 }
```

\$group

Group the documents by `total_staff_number` and by `budget` and perform summation of budgets in each group

```
db.department.aggregate([ {$group:{"_id":{"total_staff_number":"$total_staff_number",  
                                     "budget":"$budget"} } } ])
```

```
{ "_id" : { "total_staff_number" : 5, "budget" : 10000 } }  
{ "_id" : { "total_staff_number" : 25, "budget" : 100000 } }  
{ "_id" : { "total_staff_number" : 25, "budget" : 1000000 } }
```

Group the documents by `total_staff_number` and perform summation of budgets in each group

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                     "total budgets":{"$sum":"$budget"}} } ])
```

```
{ "_id" : 5, "total budgets" : 10000 }  
{ "_id" : 25, "total budgets" : 1100000 }
```

\$group

Group the documents by `total_staff_number` and find the largest budget in each group

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                "largest budget":{"$max":"$budget"}}} ])
```

Results

```
{ "_id" : 5, "largest budget" : 10000 }  
{ "_id" : 25, "largest budget" : 1000000 }
```

Group the documents by `total_staff_number` and find the smallest budget in each group

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                "smallest budget":{"$min":"$budget"}}} ])
```

Results

```
{ "_id" : 5, "largest budget" : 10000 }  
{ "_id" : 25, "largest budget" : 100000 }
```


\$group

Group the documents by `total_staff_number` and find an average budget in each group

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                "largest budget":{"avg":"$budget"}}}  ])
```

aggregate()

Other operators like `$first`, `$last` are useful when sorting is applied

What about application of aggregation operation on entire collection of documents ?

Then, (like in SQL) we assume that a collection is a single group

Count the total number of departments in a collection

```
db.department.aggregate([ {$group:{"_id":null,  
                                "total departments":{"sum:1"}}}  ])
```

aggregate()

```
{ "_id" : null, "total_departments" : 3 }
```

Results

\$group

Find the total and average budget in a collection

```
db.department.aggregate([ {$group:{"_id":null,  
                                "total budget":{"$sum:"$budget"},  
                                "average budget":{"$avg:"$budget"}}}  ])
```

aggregate()

```
{ "_id" : null, "total budget" : 1110000, "average budget" : 370000 }
```

Results

Aggregations

Outline

Aggregation framework

Operations

[\\$project](#)

[\\$match](#)

[\\$limit](#)

[\\$skip](#)

[\\$unwind](#)

[\\$group](#)

[\\$sort](#)

[\\$out](#)

[\\$count](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

27/35

\$sort

Operation `$sort` sorts the documents

Display the names and budgets of departments sorted in ascending order by budget

```
db.department.aggregate([ {$project: { "name":1, "budget":1, "_id":0}},  
                          {$sort:{"budget":1}}  ])
```

Results

```
{ "name" : "School of Astronomy", "budget" : 10000 }  
{ "name" : "School of Physics", "budget" : 100000 }  
{ "name" : "School of Computing and Information Technology", "budget" : 1000000 }
```

Display a name of a department with the largest budget, display a name of department and its budget

```
db.department.aggregate([ {$project:{"name":1, "budget":1, "_id":0}},  
                          {$sort:{"budget":-1}},  
                          {$limit:1}  ])
```

\$sort

Group the documents by `total_staff_number`, count the total number of departments in each group display the results sorted in the descending order of the total number of departments

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                "total departments":{$sum:1}}},  
                        {$sort:{"total departments":-1}} ])
```

aggregate()

```
{ "_id" : 25, "total departments" : 2 }  
{ "_id" : 5, "total departments" : 1 }
```

Results

\$sort

Group the documents by `total_staff_number`, count the total number of departments in each group display the results sorted in the descending order of the total number of departments and display only groups where total number of departments is greater than 1

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number",  
                                "total departments":{"$sum:1}}},  
                          {$sort:{"total departments":-1}},  
                          {$match:{"total departments":{"$gt:1}}}  ])
```

```
{ "_id" : 25, "total departments" : 2 }
```

Aggregations

Outline

Aggregation framework

Operations

\$project

\$match

\$limit

\$skip

\$unwind

\$group

\$sort

\$out

\$count

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

31/35

\$out

Operation `$out` saves the results of processing in a collection

Find the total number of distinct values of `total_staff_number`

<code>db.department.aggregate([{ \$group: { "_id": "\$total_staff_number" } }, { \$out: "total_distinct" }])</code>	aggregate()
<code>db.total_distinct.count()</code>	count()
<code>db.total_distinct.drop()</code>	drop()

Aggregations

Outline

Aggregation framework

Operations

\$project

\$match

\$limit

\$skip

\$unwind

\$group

\$sort

\$out

\$count

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

33/35

\$count

Operation `$count` counts the total number of documents in a pipeline

List the codes of all courses

```
db.department.aggregate([ {$unwind:"$courses"},  
                          {$project:{"code":"$courses.code", "_id":0}},  
                          {$count:"Total number of courses"} ])
```

aggregate()

Find the total number of distinct values of `total_staff_number`

```
db.department.aggregate([ {$group:{"_id":"$total_staff_number"}},  
                          {$count:"Total distinct values"} ])
```

aggregate()

References

[MongoDB Manual, Aggregation](#)

Banker K., Bakkum P., Verch S., Garret D., Hawkins T., MongoDB in Action, 2nd ed., Manning Publishers, 2016

Chodorow K. MongoDB The Definitive Guide, O'Reilly, 2013

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

35/35