

# CSCI235 Database Systems

## MongoDB Replication

Dr Janusz R. Getta

School of Computing and Information Technology -  
University of Wollongong

# Replication

## Outline

### Basics

Why replication ?

Why no-replication ?

Experiment

How does replication work ?

Commit and rollback

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

2/27

# Basics

**Replication** means automatic maintenance of data distributed over a number of MongoDB servers

**Replication** is implemented as a mechanism called **replica sets**

A **replica sets** is a group of nodes are configured such that they can to automatically synchronize their data and failover when a node disappears

Older versions of MongoDB support a method of replication called **master-slave** (now considered as deprecated) that still can be used in MongoDB v3.0

In both approaches, a single **primary node** receives all writes, and then all **secondary nodes** read and apply those writes to themselves asynchronously

**Replica sets** use the same replication mechanism as **master-slave** with additionally ensuring automated failover

# Basics

In **replica-sets** approach if the **primary node** goes offline then whenever it is possible one of the **secondary nodes** is automatically promoted to be **primary node**

Additionally, **replica sets** provide the improvements to the previous replication method like easier recovery and more sophisticated deployment topologies

In a **replica set** approach data is not considered as committed until it is written to a majority of member nodes, i.e. more than 50% of the servers

It means that if a **replica set** has only two servers then no server can be down

If the **primary node** in a **replica set** fails before it replicates its data, other members will continue accepting writes, and any not replicated data must be rolled back, meaning it can no longer be read

# Replication

## Outline

[Basics](#)

[Why replication ?](#)

[Why no-replication ?](#)

[Experiment](#)

[How does replication work ?](#)

[Commit and rollback](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

5/27

## Why replication ?

**Replication** provides safety mechanisms protecting database system from environment failures like:

- a network connection between the application and the database is lost,
- there is a loss of power,
- persistent storage device (HDD, SSD) fails

In addition to protecting against external failures, replication is important for **system durability**

When running without backup/journaling the original values of corrupted data cannot be easily restored

**Replication** can always guarantee a clean copy of the data files if a single node shuts down due to a hardware fault

**Replication** also facilitates redundancy, failover, maintenance, and load balancing

**Replication** is designed primarily for redundancy

## Why replication ?

It ensures that the contents of replicated nodes are synchronised with the primary node

The **replicas** can be located in the same place as a **primary node** or at different location connected with a **primary node** over wide area network

**Replication** is asynchronous and because of that any sort of network latency or partition between nodes have no impact on performance of the **primary node**

The modifications of the contents of replicated nodes can be delayed by a constant number of seconds, minutes, or even hours behind the primary node

It gives a chance to "move back in time" if the contents of the **primary node** is corrupted

For example, if someone accidentally drops a wrong collection of objects and a replica is "delayed in time then it is possible to restore from a **replica node**

## Why replication ?

A delayed [replica](#) gives an administrator time to react and restore data

Another application of replication is [failover](#), i.e. a situation when the [primary node](#) fails and one of the redundant nodes takes a role of the [primary node](#)

In MongoDB such "switch" is performed automatically.

Replication simplifies maintenance by allowing the high workload operations to be done on a node other than the primary at production time

For example, it is possible to run backups on a [secondary node](#) to avoid unnecessary additional load on the primary node and to avoid downtime

Replication allows for building large indexes on a [secondary node](#) simultaneously with the operations on the [primary node](#)



## Why replication ?

Then it is possible to swap the **secondary node** with the existing **primary** and then build the same index again on the new **secondary**

**Replication** allows to balance read operations across many replicas

Data can be simultaneously read from many separate replicas; it is the easiest (and a bit simplistic) way to scale up the system

# Replication

## Outline

[Basics](#)

[Why replication ?](#)

[Why no-replication ?](#)

[Experiment](#)

[How does replication work ?](#)

[Commit and rollback](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

10/27

## Why no-replication ?

**Replication** does not improve performance a lot when:

- hardware can't process the given workload; if performance becomes constrained by the number of I/O operations per second (IOPS) your disk can handle (80–100 for HDD) then reading from a replica increases your total IOPS to 100 to 200 IOPS; if writes are occurring at the same time then it consumes all IOPS ( sharding is a better option)
- the ratio of writes to reads exceeds 50%; then every write to the primary must eventually be written to all the secondary nodes as well and directing the additional reads to secondary nodes slows down replication
- an application requires consistent reads; then secondary nodes replicate asynchronously and therefore they are not guaranteed to reflect the latest writes to the primary node

# Replication

## Outline

Basics

Why replication ?

Why no-replication ?

Experiment

How does replication work ?

Commit and rollback

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

12/27

## Experiment

We use replica sets to create two replica nodes simulated by two Mongo processes at the ports 4000 and 4001

Start [Terminal](#) and process the following shell commands

```
cd ~  
mkdir node0  
mkdir node1
```

Command shell

In [Terminal](#) window start [mongod](#) server with data located in [node0](#), listening to a port 4000 and attached to a replica set "rs0"

```
mongod --dbpath node0 --port 4000 --replSet "rs0"
```

Command shell

Minimize [Terminal](#) window (do not close it !)

Open another [Terminal](#) window and start [mongod](#) server with data located in [node1](#) listening to a port 4001 and attached to the same replica set "rs0"

```
mongod --dbpath node1 --port 4001 --replSet "rs0"
```

Command shell

# Experiment

Minimize **Terminal** window (do not close it !)

Start **Terminal** and then start **mongo** client and connect to a server listening at a port **4000**

```
mongo --port 4000
```

Command shell

At **mongo** client prompt **>** process the following commands

```
rs.initiate()
```

Initiate()

Reply is such that the current node is **SECONDARY**

```
rs.conf()
```

Initiate()

The current node became **PRIMARY**

Add **mongod** server at **4001** as a secondary node

```
rs.add("localhost:4001")
```

add()

# Experiment

Check status

```
rs.status()
```

```
status()
```

Open yet another [Terminal](#) window and connect to [mongod](#) server at 4001

```
mongo --port 4001
```

```
Command shell
```

Make the node slave

```
rs.slaveOk()
```

```
slaveOk()
```

## Experiment

Now, to test replication set we shall create a new collection and we shall insert a new document at mongo server at a port 4000; replication mechanism supposed to copy insertion to mongod server at a port 4001

Return to a window where mongo client is connected to mongod server at a port 4000 and process the following commands

```
use mydb
db.names.insert({"full-name": "James Bond"})
db.names.find()
```

port 4000

Move to a window where mongo client is connected to mongod server at a port 4001 and execute the commands

```
use mydb
db.names.find()
```

port 4001



## Experiment

An attempt to insert a new document by a client connected to **mongod** server at a port **4001** fails due to slave mode of the server

To shutdown both replication servers execute in both windows the following statements

```
use admin  
db.shutdownServer()
```

shutdown

After replication set has been created and after shutdown it is possible to restart it in the following way

In **Terminal** window start **mongod** server with data located in **node0**, listening to a port **4000** and attached to a replica set "**rs0**"

```
mongod --dbpath node0 --port 4000 --replSet "rs0"
```

Command shell

## Experiment

Minimize [Terminal](#) window (do not close it !)

Open another [Terminal](#) window and start mongod server with data located in [node1](#) listening to a port [4001](#) and attached to the same replica set

```
mongod --dbpath node1 --port 4001 --replSet "rs0"
```

[Command shell](#)

Minimize [Terminal](#) window (do not close it !)

Open yet another [Terminal](#) window and connect to [mongod](#) server at [4001](#)

```
mongo --port 4001
```

[Command shell](#)

Make the node slave

```
rs.slaveOk()
```

[slaveOk\(\)](#)

# Experiment

Start [Terminal](#) and then start [mongo](#) client and connect to a server listening at a port [4000](#)

```
mongo --port 4000
```

Command shell

Return to a window where [mongo](#) client is connected to [mongod](#) server at a port [4000](#) and process the following commands

```
use mydb  
db.names.insert({"full-name": "Harry Potter"})  
db.names.find()
```

port 4000

Move to a window where [mongo](#) client is connected to [mongod](#) server at a port [4001](#) and execute the commands

```
use mydb  
db.names.find()
```

port 4001

# Replication

## Outline

[Basics](#)

[Why replication ?](#)

[Why no-replication ?](#)

[Experiment](#)

[How does replication work ?](#)

[Commit and rollback](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

20/27

## How does replication work ?

Replica sets rely on two basic mechanisms: an **oplog** and a **heartbeat**

**Oplog (operation log)** is a space restricted collection that lives in a database called **local** on every replica node and records all changes to the data

Every time a client writes to the primary node, an entry with enough information to reproduce the write is automatically added to the primary node's **oplog**.

Once the write is replicated to a given secondary node then its **oplog** also stores a record of the write

When a given secondary node is ready to update itself, it performs the following actions:

- first, it looks at the timestamp of the latest entry in its own **oplog**
- next, it queries the primary node's **oplog** for all entries greater than that timestamp
- finally, it writes the data and adds each of those entries to its own **oplog**

## How does replication work ?

Then in case of failover, any secondary promoted to primary will have an **oplog** that the other secondaries can replicate from; it enables replica set recovery

Secondary nodes use **long polling** to immediately apply new entries from the primary's oplog

**Long polling** means the secondary makes a long-lived request to the primary

When the primary receives a modification, it responds to the waiting request immediately

The secondary nodes will usually be almost completely up to date

If the secondary nodes fall behind because of network partitions or maintenance on secondaries, the latest timestamp in each secondary's oplog can be used to monitor any replication lag

The replica set **heartbeat** facilitates election and failover

## How does replication work ?

By default, each replica set member pings all the other members every two seconds

As long as every node remains healthy and responsive, the replica set continue its work

Every replica set wants to ensure that exactly one primary node exists at all times

With no majority, the primary demotes itself to a secondary

If the heartbeats fail due to some kind of network partition, the other nodes will still be online

If the arbiter and secondary are still up and able to see each other, then according to the rule of the majority, the remaining secondary will become a primary

# Replication

## Outline

[Basics](#)

[Why replication ?](#)

[Why no-replication ?](#)

[Experiment](#)

[How does replication work ?](#)

[Commit and rollback](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

24/27



## Commit and rollback

Writes are not considered as committed until they are replicated to a majority of nodes

Operations on a single document are always atomic and operations that involve multiple documents are not atomic

Consider the following scenario:

- a series of writes to the primary node did not get replicated to the secondary node,
- primary node goes offline and the secondary is promoted to primary and new writes go to primary
- old primary comes back online as secondary and tries to replicate from the new primary
- old primary has a series of writes that don't exist in the new primary node **oplog**
- It triggers a **rollback**

**Rollback** reverses all writes that were never replicated to a majority of nodes

## Commit and rollback

The writes are removed from both the secondary's oplog and the collection where they reside

If a secondary node has registered a deleted document, the node will look for the deleted document in another replica and restore it to itself

The same is true for dropped collections and updated documents

The reverted writes are stored in the rollback subdirectory of the path in the relevant node

For each collection with rolled-back writes, a separate BSON file will be created the filename of which includes the time of the rollback

In the event that you need to restore the reverted documents, you can examine these BSON files using the `bsondump` utility and manually restore them, possibly using `mongorestore`

# References

Banker K., Bakkum P., Verch S., Garret D., Hawkins T., MongoDB in Action, 2nd ed., Manning Publishers, 2016

[MongoDB Manual, Replication](#)

[TOP](#)

Created by Janusz R. Getta, CSCI235 Database Systems, Spring 2020

27/27