

CSCI235 Database Systems

NoSQL, NewSQL, and Other Database Systems

Dr Janusz R. Getta

School of Computing and Information Technology -
University of Wollongong

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

XML Database Systems

The **Extensible Markup Language (XML)** was defined by the **WWW Consortium (W3C)** as a document markup language in late 1990s

XML has become a major format for data exchange, flexibly structured documents, or configuration files

XML document consists of nested **XML elements**

XML element consists of **start tag**, **data**, and **end tag**

XML element may contain other **XML elements**

XML element can be described by the attributes located within **start tag**

Attributes provide a meta description or identification of XML element

XML document may contain additional information like **version declarations**, **entity declarations**, **comments**, and **processing instructions**

XML namespaces (XMLNS) are useful to have unique naming of elements

Namespace is usually **Uniform Resource Identifier (URI)** which denotes scope and origin of the document

XML Database Systems

Sample XML document

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE booking SYSTEM "/xml-resources/dtd/booking.dtd">
<booking>
  <hotel ID="Hongkong Sheraton">
    <rooms>
      <room>
        <category>suite</category>
        <price>1234</price>
      </room>
      <room>
        <category IDREF="Sydney Intercontinental">standard</category>
      </room>
    </rooms>
  </hotel>
</booking>
```

XML

XML Database Systems

Document Type Definition (DTD) is the simplest way to specify a schema for **XML documents**

XML Schema is more sophisticated schema language than DTD

XML Schema allows for data typing withing a system defined types, creating user defined types, and inheritance

XML Schema supports constraints on the number of occurrences of sublements, on maximum, minimum values, and an advanced ID referencing

XML has 3 query/transformation languages:

- **XPath**: A language for navigation in XML document
- **XQuery**: A complex query and transformation language that uses XPath as navigational tool
- **XSLT**: A query and transformation language

XML Database Systems

Native XML database systems: XML document management systems

- **eXistDB**: Java-based open source XML database; supports indexing, XPath, XQuery, XML:DBI (Java API)
- **BaseX**: Java-based system that stores XML documents in a tabular representation; support indexing, conversion to JSON, XML:DB API

XML support in relational a database systems: columns of type **XMLType**, e.g. ANSI SQL definition (**SQL/XML**) and Oracle, Postgres, DB2, MySQL, and SQL Server

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Key-value Stores and Document Database Systems

Key-value pair is a tuple of two strings ("key" , "value")

A **key-value store** stores key-value pairs

A value is retrieved by specifying a key

A **key-value store** is a prototype of a **schemaless** database system

It is possible to put any **key-value** into a store and no restrictions are enforced on a format or structure of a value

A **key-value store** offers three operations: **put**, **get**, and **delete**

A value in **key-value** pair may contain other **key-value** pairs or sequences of keys

Key-value Stores and Document Database Systems

Dynamo DB is a **key-value database** implemented by Amazon

Architectural characteristics:

- The system uses a hash value of a key to determine the nodes in a cluster responsible for the key; nodes can be added and removed with minimal rebalancing overhead
- To determine one of n nodes in cluster where a key-value is saved $\text{hash}(\text{key})$ is computed modulo n
- Addition of a new node is performed by re-hashing
- An application can determine trade-offs between consistency, read performance and write performance
- It is possible to determine **strong consistency**, **eventual consistency**, and **weak consistency**
- Since write operations are never blocked it is possible that there will be multiple versions of an object in the system
- Versions can be merged by a data store itself or sometimes need to be resolved by the application or user
- For instance, if the buyer updates his or her shopping cart from two computers, the resulting cart may have duplicate items that he or she may need to remove

Key-value Stores and Document Database Systems

A **document database** stores and retrieves documents, which can be in JSON, BSON, etc format

Documents are self-describing, hierarchical data structures

The structures of documents are similar to each other but not exactly the same

Documents are stored as the value part of **key-value** store

```
{
  "firstname": "James",
  "likes": ["biking", "hiking", "viking"],
  "city": "Boston",
  "friend": null,
  "addresses": [{ "state": "NSW", "city": "Sydney" },
                { "state": "Vic", "city": "Melbourne" } ]
}
```

JSON

Key-value Stores and Document Database Systems

Other **key-value** stores and document database systems:

- **Riak**: it groups key-value pairs called as Riak objects into logical units called buckets; buckets can be configured by defining bucket type (specific configuration for a bucket); the system comes with several options to configure a storage backend
- **Redis**: it is an advanced **in-memory key-value** store with command line interface; it supports several data types and data structures like strings, linked lists, unsorted sets, sorted sets, hash maps, bit arrays
- **MongoDB**: it is a document database with **BSON** as its storage format
- **CouchDB**: it is an **Erlang** based document database; it stores **JSON** documents in so called databases and it supports multi-version concurrency control; it exposes conflict handling to the users: a user has to submit the most recent revision number of a document when writing data into it
- **Couchbase**: it is a document database that merges the features of **CouchDB** and **Membase**; it stores **JSON** documents and other formats like binary and serialized data in buckets; each document is stored under a unique key; it implements multiversion concurrency control
- Other other systems: **ArangoDB**, **OrientDB**, **RavenDB**, **RethinkDB**

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Graph databases

Graph databases store entities and relationships between the entities

Nodes are instances of objects in the applications

Edges have properties and have directional significance

```
Node james = graphDb.createNode();
james.setProperty("name", "James");
Node harry = graphDb.createNode();
harry.setProperty("name", "Harry");
james.createRelationship(harry, FRIEND);
harry.createRelationship(james, FRIEND);
```

Graph data model

Graph databases ensure consistency through **ACID**-compliant transactions

Graph databases use domain specific languages for traversing graph structures

Properties of nodes can be indexed

A commonly used technique for horizontal scaling is **sharding**

Any link-rich domain is well suited for **graph databases**

Graph databases

Implementation of logical graph data model in a relational database system requires indexing to speed up processing of join operations (**index lookups**)

Efficient real-time graph processing requires a way to move through the graph structure without index lookups (**index-free adjacency**)

In **index-free adjacency** implementation of graph data model each node knows the physical location of all adjacent nodes

There is no need to use indexes to efficiently navigate the graph, because each node “points” to all adjacent nodes

It creates the problems with distribution of data among many servers

Graph compute engines enable graph processing that needs to work across distributed datasets

Graph databases

Neo4j is a Java-based graph database based on **Property Graph** data model

Property Graph data model provides a richer model for representing complex data than **RDF** (presented later) by associating both nodes and relationships with attributes

Neo4j supports billions of nodes, **ACID** compliant transactions, and multiversion consistency

Neo4j implements a declarative graph query language **Cypher** comparable to **SQL** or **SPARQL** (query language for **RDF**)

Graph databases

The following sequence of statements creates the data nodes and edges in **Neo4j** database

```
CREATE (TheMatrix:Movie {title:'The Matrix',  
                        released:1999,  
                        tagline:'Welcome to the Real World'})
```

Neo4j

```
CREATE (JohnWick:Movie {title:'John Wick',  
                      released:2014,  
                      tagline:'Silliest Keanu movie ever'})
```

Neo4j

```
CREATE (Keanu:Person {name:'Keanu Reeves', born:1964})  
CREATE (AndyW:Person {name:'Andy Wachowski', born:1967})
```

Neo4j

```
CREATE  
(Keanu)-[:ACTED_IN {roles:['Neo']}]> (TheMatrix),  
(Keanu)-[:ACTED_IN {roles:['John Wick']}]> (JohnWick),  
(AndyW)-[:DIRECTED]> (TheMatrix)
```

Neo4j

Graph databases

The following queries retrieve information from **Neo4j** database

```
MATCH (kenau:Person {name:"Keanu Reeves"})  
RETURN kenau;
```

Neo4j

```
MATCH (kenau:Person {name:"Keanu Reeves"})  
-[:ACTED_IN]-  
(movie)<-[:ACTED_IN]-(coStar)  
RETURN coStar.name;
```

Neo4j

```
MATCH (kenau:Person {name:"Keanu Reeves"})  
-[*1..2]-  
(related)  
RETURN distinct related;
```

Neo4j

Graph databases

Resource Description Framework (RDF) (late 1990) is a Web standard developed for modeling of Web resources and the relationships between them

Information in **RDF** is expressed as triples **subject:predicate:object**

```
TheMatrix:is:Movie  
Kenau:is:Person  
Kenau:StarredIn:TheMatrix
```

RDF

RDF was intended as the means of creating a formal database of resources on the web together with the dependencies that these resources relied on

RDF graphs can be stored in a variety of formats including XML, or even within tables inside a relational database

A widely used query language for **RDF** graphs is **SPARQL** that has an **SQL**-like syntax

Graph databases

The other most popular systems include:

- **InfiniteGraph**: it is an enterprise distributed graph database implemented in Java and build on former object-oriented database **ObjectivityDB**
- **OrientDB**: it is the first multi-model NoSQL database system with a graph database engine; it provides graph, document, key-value, and object view of data
- **FlockDB**: it is an open source distributed, fault-tolerant graph database for managing wide but shallow network graphs; it is not designed for multi-edge graph traversal but rather for rapid set operations on graph connections
- **AllegroGraph**: it is a closed source triple store which is designed to store **RDF** triples; it supports transactions, automatic indexing, querying with SPARQL and reasoning with RDFS++ and Prolog
- **Ontotext**: it provides **RDF** view of data for content management, knowledge discovery and semantic search
- **GraphDB**: it is a **semantic graph database**, compliant with W3C Standards; it provides a view of a database as **RDF** triplestores

Graph databases

The other most popular systems include:

- **StarDog**: it allows to query, search, and analyze enterprise data using **knowledge graph** technology; it supports unification with relational database systems, **ACID** transaction processing, query answering for location aware (geospacial) graphs
- **Oracle Spatial and Graph**: it includes high performance, enterprise-scale, commercial spatial and graph database and analytics for Oracle Database 12 it supports large-scale Geographic Information Systems (**GIS**)
- **Apache Giraph**: it is an iterative graph processing framework, built on top of **Apache Hadoop**;
- **GraphX**: it unifies Extract Transform and Load (**ETL**), performs exploratory analysis, and iterative graph computation within a single system; it provides a view of the same data as both graphs and collections

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Column stores

Column store (or columnar database management system) is a database management system (**DBMS**) that stores data tables by column rather than by row

Practical use of a column store versus a row store differs little in the relational DBMS world

To reduce overhead imposed DML operations columnar databases generally implement some form of **write-optimized delta store**

Delta store is optimized for frequent writes like tables stored in row-by-row format

Delta store is merged with the main columnar-oriented store

The merge will occur periodically, or whenever the amount of data in the delta store exceeds a threshold

Prior to the merge, queries access both the delta store and the column store in order to return complete and accurate results

Column stores

KDB was the first commercially available column-oriented database developed in 1993 followed in 1995 by **Sybase IQ**

Sybase IQ (mid 1990) ("Intelligent Query") data warehousing platform implemented relational tables in column-by-column mode

C-Store and later on **Vertica** (2005-2011) implemented a native columnar database

MonetDB was released under an open-source license in 2004

Vertica was eventually developed out of **C-Store**, while the **MonetDB** related project evolved into VectorWise

Druid is a column-oriented data store that was open-sourced in late 2012 and now it is used by numerous organizations

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Extensible Record Stores

Extensible Record Store organizes data into **extensible tables**

Extensible table consists of **rows**

Each **row** is uniquely identified by a **row key**

Data within a **row** is grouped by a **column family**

Column families have an important impact on the **physical implementation** of **extensible table**

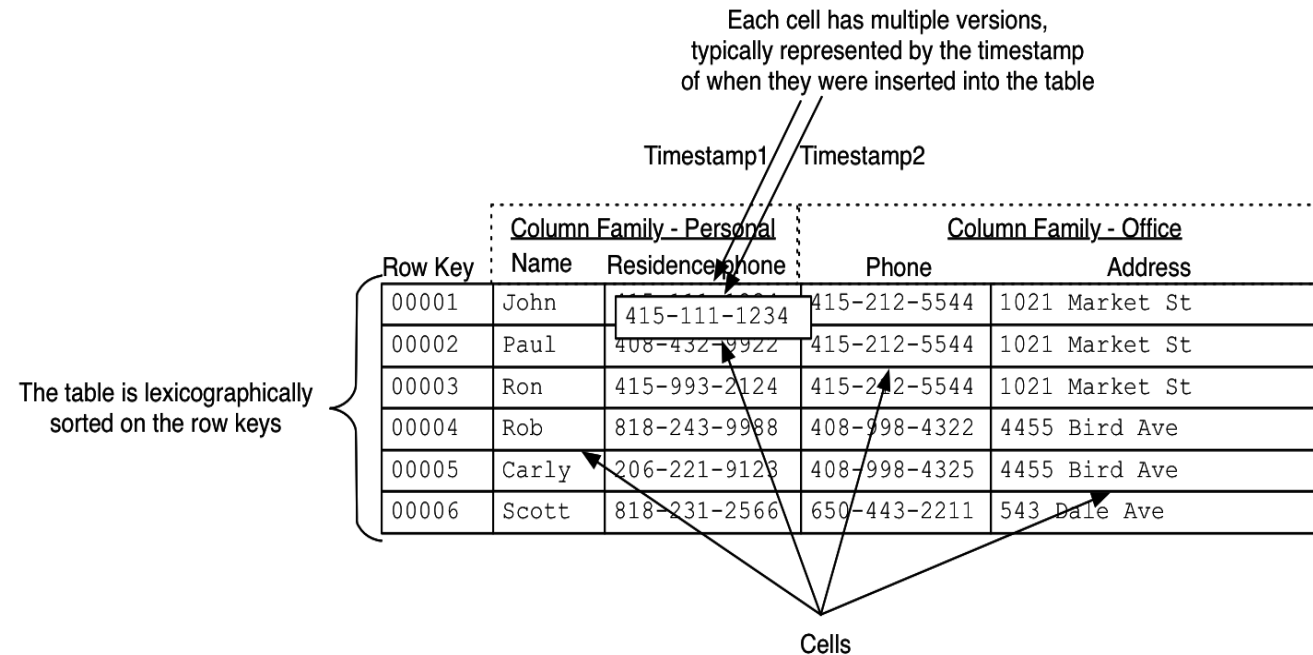
Every **row** has the same **column families** although some of them can be **empty**

Data within a **column family** is addressed via its **column qualifier**, or simply, **column name**

Hence, a combination of **row key**, **column family**, and **column qualifier** uniquely identifies a **cell**

Values in cells do not have a data type and are always treated as **sequences of bytes**

Extensible Record Stores



Values within a cell have multiple versions

Versions are identified by their version number, which by default is a timestamp when the cell was written

Extensible Record Stores

If a **timestamp** is not determined at write time, then the **current timestamp** is used

If a **timestamp** is not determined during a read, the **latest one** is returned

The **maximum allowed number of cell value versions** is determined for each **column family**

Foundation of **Extensible Record Stores** is Google's system called **Big Table**

The other systems include

- **Apache Cassandra**: it stores column families in so called "keyspace"; create table command creates a new column family in the keyspace; indexes can be created on columns
- **Apache HBase**: it stores tables in namespaces; it offers command line interface, SQL interface called **Phoenix**; and **Java API**
- **Hypertable**: it stores tables in namespaces; it provides a query language **HQL** similar to **SQL**
- **Apache Accumulo**: it also stores tables in namespaces and it also provides a command line interface and **Java API**

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Object-Oriented and Object-Relational Database Systems

Object-Oriented database systems have its roots in a concept of **persistent programming language** that distinguishes between transient (RAM memory) data structures and persistent (HDD, SSD memory) data structures

Object-Oriented database systems allow to store and to operate on both persistent and transient objects

Object Definition Language (ODL) can be used to define classes of objects, properties, associations, generalization and aggregation hierarchies, and methods

Object Query Language (OQL) provides SQL-like syntax to express queries on instances of objects in a database

One of the approaches to implementation of persistently stored objects was to use conventional relational database systems as an underlying storage engine

Such approach required **Object-Relational** mapping of object data model into tabular data model

Object-Oriented and Object-Relational Database Systems

Object-Relational mapping provides a view of a database as a collection of tables filled with objects instead of rows and it is called as **Object-Relational** data model

Object-Relational database system a database is a collection of tables (containers) filled with objects that contains references to other objects in the same or in te other tables (containers)

Object-Relational Mapping (ORM) tools are avaible for a wide range of **Object-Oriented** programming languages

The standards for **Java** object persistence are **Java Persistence API (JPA)** and **Java Data Objects (JDO)** API

Pure **Object-Oriented** database systems are almost non existent any more unless we consider **JSON** data model based systems as partial implementation of **Object-Oriented** database systems

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

In-Memory Database Systems

Solid state disk (SSD) technology has allowed for implementation of in-memory database systems

SSDs store bits of information in cells

A single-level cell (SLC) contains one bit of information per cell, while a multi-level cell (MLC) SSD contains usually only two but sometimes three bits per each cell

Cells are arranged in pages of about 4K and pages are arranged in blocks that typically contain 256 pages

Read operations, and initial write operations, require only a single-page IO

Changing the contents of a page requires an erase and overwrite of a complete block, i.e. it is two times slower than read

Many traditional relational databases perform relatively poorly on SSDs

This is because critical IO has been isolated to sequential write operations for which hard disk drives perform well, but that represent the worst possible workload for SSD

In-Memory Database Systems

When a transaction commits in an **ACID-compliant** relational database, the transaction record is usually written immediately to a sequential transaction log (called a redo log in some databases)

When the log is moved to **SSD** performance does not improve because of lesser **SSD** write performance

SSD based systems (**Aerospike**) use main memory to store indexes to the data while keeping the data always on flash

A more paradigm-shifting trend has been the increasing practicality of storing complete databases in **main memory** (RAM)

Architectural problems of main memory databases:

- **Cache-less architecture**: there is no point caching in memory what is already stored in memory
- **Alternative persistence model**: data in memory disappears when the power is turned off

In-Memory Database Systems

In-memory databases either

- replicate data to other members of a cluster or
- write complete database images to disk files or
- write out transaction/operation records to an append-only disk file (called a transaction log or journal)

The most popular in-memory systems include:

- **TimesTen**: (owned by Oracle since 2005); it is entirely based on transactional relational database architecture; it can be used as a standalone in-memory database or as a caching database supplementing the traditional disk-based Oracle RDBMS; when the database is started, all data is loaded from checkpoint files into main memory; periodically or when required database data is written to checkpoint files
- **Redis**: (owned by VMware since 2013); it is an in-memory key-value store; it uses a [snapshot files](#) to store the copies of the entire system at a given point in time; [append only file](#) keeps a journal of changes that can be used to “roll forward” the database from a snapshot in the event of a failure; configuration options allow the users to configure writes to the [append files](#) after every operation, at one-second intervals, or based on operating-system-determined flush intervals

In-Memory Database Systems

The most popular in-memory systems include:

- **SAP HANA**: (owned by SAP, 2010) it is a relational database that combines in-memory technology with a columnar storage option, installed on an optimized hardware configuration; relational tables in **HANA** can be configured for row-oriented or columnar storage; on start-up, row-based tables and selected column tables are loaded from the database files and other column tables can be loaded on demand; reads and writes to row-oriented tables are applied directly; updates to column-based tables are buffered in the delta store; queries against column tables must read from both the delta store and the main column store
- **VoltDB**: it is a **hybrid** (in-memory/HDD) ACID-compliant relational database based on **shared nothing** (**sharding**) architecture; it supports SQL access from within pre-compiled Java stored procedures; it relies on horizontal partitioning down to the individual hardware thread to scale
- **DB2 with BLU acceleration**: (owned by IBM) it is a hybrid relational database that stores frequently used data on **SSD**; it allows for better performance at a price closer to **HDD** based systems than purely **SDD** based systems, e.g. Oracle's Exadata database machine

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Data Stream Processing Systems

A **data stream** is an infinite sequence of transient values and when recorded persistent values

Data streams can be produced by sensor networks or by network traffic monitoring, stock exchange monitoring, etc

Data Stream Management System (DSMS) processes data streams by interminably running **continuous queries** over and over again

DSMS must handle the queries that might be running for days, months, or even years

DSMS must aggregate data from the entire stream, e.g. overall average of all values, or from a subset of data called as sliding window, e.g. an average of values in the last 30 days

Majority of **DSMS** are based on the relational data model and their continuous query languages are similar to SQL

In such a case data items from the streams is represented as a pair **(timestamp,tuple)** where tuple corresponds to a row in an infinite relational table

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

Array Database Systems

Array Database Systems organize data along multiple dimensions and can be used to store and to manipulate data with complex structures, e.g. astronomical data obtained from satellite observations

In the **Array Data Model** data are stored in **multidimensional arrays**

Each **cell** in an array contains a **tuple** of certain length

The elements of such **tuple** can be either **atomic values** or **arrays** by themselves

Array Data Model can express an arbitrary number of **array nestings**

Tuples are addressed by specifying the corresponding **dimensions**

In an **Array Database System SciDB** individual scalar values can be addressed by named attributes

Array Database Systems offer specialized functions to manipulate array data, like different types of aggregations, operation on matrices, and even joins

NoSQL, NewSQL, and Other Database Systems

Outline

[XML Database Systems](#)

[Key-value Stores and Document Database Systems](#)

[Graph Database Systems](#)

[Column Stores](#)

[Extensible Record Stores](#)

[Object-Oriented and Object-Relational Database Systems](#)

[In-Memory Database Systems](#)

[Data Stream Processing Systems](#)

[Array Database Systems](#)

[NewSQL Database Systems](#)

NewSQL Database Systems

A term **NewSQL** describes adoption of the design principles underlying **NoSQL** systems to build new distributed relational database systems with SQL interface

Such re-design of conventional relational DBMSs would make them **scalable**, i.e. it would be possible to store relational data in a number of server while efficiently processing SQL queries

Additionally, **NewSQL** systems suppose to transparently handle hardware failures in a network while still maintaining **ACID** compliant transaction protocols

The first type of **NewSQL** systems are completely new database platforms designed to operate in a distributed cluster of shared-nothing nodes, in which each node owns a subset of the data, e.g. **Google Spanner**, **CockroachDB**, **Clustrix**, **VoltDB**, **MemSQL**, **NuoDB**, and **Trafodion**

The second type are highly optimized storage engines for SQL that provide the same programming interface as SQL, but scale better than built-in engines, e.g. **MySQL Cluster**, **Infobright**, **TokuDB**, and **MyRocks**

References

Wiese L. Advanced Data Management for SQL, NoSQL, Cloud and Distributed Databases, De Gruyter Oldenburg, 2015

Harrison G., Next Generation Databases NoSQL NewSQL Big Data, Apress, 2015